

Attention Mechanism

The attention mechanism allows a model to focus on relevant parts of the input sequence while processing information. In self-attention, each token interacts with every other token through Query, Key, and Value vectors. Attention is a core component of Transformer architectures and powers modern Large Language Models. Applications include machine translation, summarization, question answering, and chatbots.